CSCI461 Big Data Assignment #1

Overview

This project processes a dataset using a pipeline built with Docker and Python. It includes steps for data loading, preprocessing, exploratory data analysis (EDA), visualization, and K-means clustering. The output files are stored in the res directory.

Requirements

- Docker
- Dataset (iris.csv used in this example)

Project Structure

- **Dockerfile**: Defines the Docker container setup with necessary packages.
- **load.py**: Loads the dataset and saves it as loaded_data.csv.
- **dpre.py**: Performs data cleaning, transformation, reduction, and discretization on the dataset. The output is saved as res_dpre.csv.
- **eda.py**: Generates insights from the data and saves them in eda-in-1.txt, eda-in-2.txt, and eda-in-3.txt.
- **vis.py**: Creates a visualization and saves it as vis.png.
- **model.py**: Applies K-means clustering with k=3 and saves the cluster counts in k.txt.

Setup and Execution

1. Build the Docker Image

In the bd-a1 directory (where the Dockerfile is located), build the Docker image:

bash

Copy code

docker build -t bd-a1-image .

2. Run the Docker Container

Run the container interactively to access the bash shell:

bash

Copy code

docker run -it --name bd-a1-container bd-a1-image

3. Execute the Pipeline

Inside the Docker container, navigate to /home/doc-bd-a1/ and execute each Python script in the following order:

1. **Load the Data**:

bash

Copy code

python3 load.py /home/doc-bd-a1/iris.csv

2. **Preprocess the Data**:

bash

Copy code

python3 dpre.py

3. **Perform EDA**:

bash

Copy code

python3 eda.py

4. **Generate Visualization**:

bash

Copy code

python3 vis.py

5. **Apply K-means Clustering**:

bash

Copy code

python3 model.py

4. Copy Output Files to the Local Machine

After executing the pipeline, copy the generated files from the container to the res directory on your local machine:

bash

Copy code

docker cp bd-a1-container:/home/doc-bd-a1/res_dpre.csv C:\Users\Salma\Desktop\bd-a1\res\

docker cp bd-a1-container:/home/doc-bd-a1/eda-in-1.txt C:\Users\Salma\Desktop\bd-a1\res\

docker cp bd-a1-container:/home/doc-bd-a1/eda-in-2.txt C:\Users\Salma\Desktop\bd-a1\res\

docker cp bd-a1-container:/home/doc-bd-a1/eda-in-3.txt C:\Users\Salma\Desktop\bd-a1\res\

docker cp bd-a1-container:/home/doc-bd-a1/vis.png C:\Users\Salma\Desktop\bd-a1\res\

docker cp bd-a1-container:/home/doc-bd-a1/k.txt C:\Users\Salma\Desktop\bd-a1\res\

Output Files

- **res_dpre.csv**: Preprocessed data.
- **eda-in-1.txt, eda-in-2.txt, eda-in-3.txt**: EDA insights.
- **vis.png**: Visualization image.
- **k.txt**: Cluster counts from K-means.

Troubleshooting

- Ensure that all Python scripts are copied into the /home/doc-bd-a1/ directory in the container before execution.
- If you encounter any issues with file paths, ensure they are specified in Unix-style (e.g., /home/doc-bd-a1/iris.csv).

Bonus (Optional)

1. **Push Docker Image to Docker Hub**:
    - Tag and push the Docker image:

bash

Copy code

docker tag bd-a1-image salmaheshamsalem123/bd-a1-image

docker push salmaheshamsalem123/bd-a1-image

2. **Push Project to GitHub**:
    - Create a repository on GitHub, add your files, commit, and push.